

```
In [ ]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
import statsmodels.api as sm
import statsmodels.formula.api as smf
from sklearn.linear_model import LinearRegression
sns.set()
```

```
In [ ]: df = pd.read_csv("Auto.csv")

#remove any non-numerical data (? values)
df[df.columns[:-2]] = df[df.columns[:-2]].apply(pd.to_numeric, errors='coerce')
df = df.dropna()
df = df.reset_index(drop=True)
```

QUESTION 1:

a)

```
In [ ]: x = df['horsepower'].values.reshape(-1, 1)
y = df['mpg'].values.reshape(-1, 1)
reg = LinearRegression().fit(x, y)
# print("Coeff of determination (R^2): {}".format(reg.score(x, y)))
X = sm.add_constant(x)
model = sm.OLS(y,X).fit()
print(model.summary())

print("Predicted mpg with a horsepower of 95: {}".format(reg.predict([[95]][0][0])))
```

OLS Regression Results

=====					
Dep. Variable:	y	R-squared:	0.606		
Model:	OLS	Adj. R-squared:	0.605		
Method:	Least Squares	F-statistic:	599.7		
Date:	Tue, 15 Nov 2022	Prob (F-statistic):	7.03e-81		
Time:	18:40:06	Log-Likelihood:	-1178.7		
No. Observations:	392	AIC:	2361.		
Df Residuals:	390	BIC:	2369.		
Df Model:	1				
Covariance Type:	nonrobust				
=====					
	coef	std err	t	P> t	[0.025 0.975]

const	39.9359	0.717	55.660	0.000	38.525 41.347
x1	-0.1578	0.006	-24.489	0.000	-0.171 -0.145
=====					
Omnibus:	16.432	Durbin-Watson:	0.920		
Prob(Omnibus):	0.000	Jarque-Bera (JB):	17.305		
Skew:	0.492	Prob(JB):	0.000175		
Kurtosis:	3.299	Cond. No.	322.		
=====					

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Predicted mpg with a horsepower of 95: 24.94061135257337

i) As the coeff of determination (R^2) is relatively close to 1, there is a relationship between the predictor and the response. The p-value is also very small for the linear regression with mpg and horsepower


ii) The relationship between the predictor and response is strong. The p value is small when performing a t-test, and the coefficient of determination is also .606, which indicates that almost 60.6% of the variability in mpg can be explained by horsepower.

iii) The relationship between the predictor and response is negative. We can see this, as the coefficient of x1 (our predictor) is negative.

iv) The predicted mpg with a horsepower of 95 is 24.94061135257337

b)

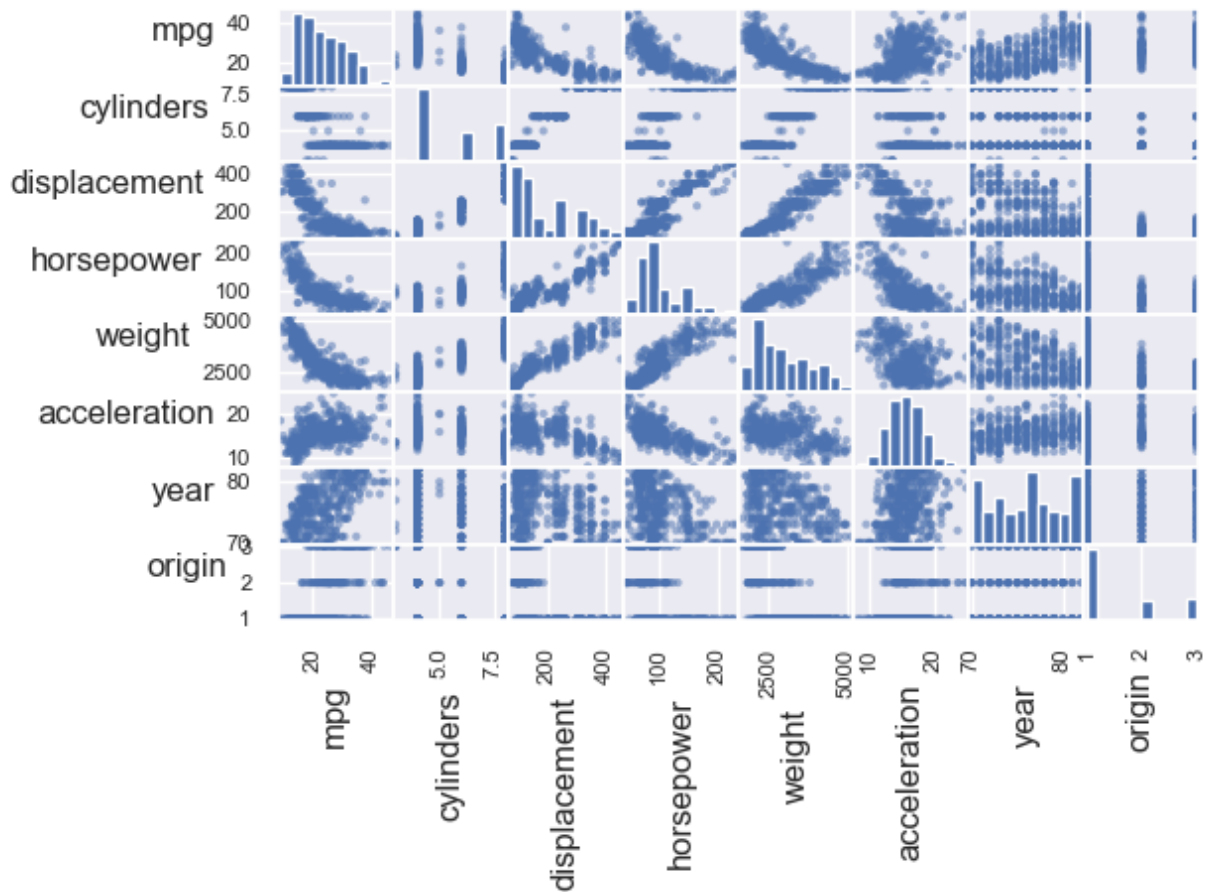
```
In [ ]: ypred = reg.predict(x)
df.plot(x='horsepower', y='mpg', style='o')
plt.title("horsepower vs mpg")
plt.ylabel('mpg')
plt.plot(x, ypred)
plt.show()
```



2.

a)

```
In [ ]: axes = pd.plotting.scatter_matrix(df)
for ax in axes.flatten():
    ax.xaxis.label.set_rotation(90)
    ax.yaxis.label.set_rotation(0)
    ax.yaxis.label.set_ha('right')
plt.tight_layout()
plt.gcf().subplots_adjust(wspace=0, hspace=0)
plt.show()
```



b)

```
In [ ]: corrM = df.corr() #no names
print(corrM)
```

	mpg	cylinders	displacement	horsepower	weight	\
mpg	1.000000	-0.777618	-0.805127	-0.778427	-0.832244	
cylinders	-0.777618	1.000000	0.950823	0.842983	0.897527	
displacement	-0.805127	0.950823	1.000000	0.897257	0.932994	
horsepower	-0.778427	0.842983	0.897257	1.000000	0.864538	
weight	-0.832244	0.897527	0.932994	0.864538	1.000000	
acceleration	0.423329	-0.504683	-0.543800	-0.689196	-0.416839	
year	0.580541	-0.345647	-0.369855	-0.416361	-0.309120	
origin	0.565209	-0.568932	-0.614535	-0.455171	-0.585005	

	acceleration	year	origin
mpg	0.423329	0.580541	0.565209
cylinders	-0.504683	-0.345647	-0.568932
displacement	-0.543800	-0.369855	-0.614535
horsepower	-0.689196	-0.416361	-0.455171
weight	-0.416839	-0.309120	-0.585005
acceleration	1.000000	0.290316	0.212746
year	0.290316	1.000000	0.181528
origin	0.212746	0.181528	1.000000

C:\Users\Bernhard\AppData\Local\Temp\ipykernel_41592\2646264494.py:1: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.

```
corrM = df.corr() #no names
```

c)

```
In [ ]: dfplot = df.drop(["mpg", "name"], axis = 1)

x = dfplot
y = df['mpg'].values.reshape(-1, 1)
reg = LinearRegression().fit(x, y)
print("Coeff of determination (R^2): {}".format(reg.score(x, y)))
X = sm.add_constant(x)
model = sm.OLS(y, X).fit()
print(model.summary())
```

Coeff of determination (R^2): 0.8214780764810599

OLS Regression Results

```
=====
Dep. Variable:          y      R-squared:          0.821
Model:                  OLS    Adj. R-squared:      0.818
Method:                 Least Squares    F-statistic:      252.4
Date:                   Tue, 15 Nov 2022    Prob (F-statistic): 2.04e-139
Time:                   18:40:08    Log-Likelihood:    -1023.5
No. Observations:      392    AIC:              2063.
Df Residuals:          384    BIC:              2095.
Df Model:              7
Covariance Type:       nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
const	-17.2184	4.644	-3.707	0.000	-26.350	-8.087
cylinders	-0.4934	0.323	-1.526	0.128	-1.129	0.142
displacement	0.0199	0.008	2.647	0.008	0.005	0.035
horsepower	-0.0170	0.014	-1.230	0.220	-0.044	0.010
weight	-0.0065	0.001	-9.929	0.000	-0.008	-0.005
acceleration	0.0806	0.099	0.815	0.415	-0.114	0.275
year	0.7508	0.051	14.729	0.000	0.651	0.851
origin	1.4261	0.278	5.127	0.000	0.879	1.973

```
=====
Omnibus:                 31.906    Durbin-Watson:          1.309
Prob(Omnibus):           0.000    Jarque-Bera (JB):        53.100
Skew:                    0.529    Prob(JB):                2.95e-12
Kurtosis:                4.460    Cond. No.                8.59e+04
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 8.59e+04. This might indicate that there are strong multicollinearity or other numerical problems.

i) there is a strong relationship between the predictors and the response

ii) weight, year, origin, and displacement seem to have statistically significant relationships to the response.

iii) the coefficient for 'year' suggests that as the year increases, the mpg of the vehicle increases as well by a factor of .75 assuming that all other predictors are constant.

d)



```
In [ ]: # ypred = reg.predict(x)
# df.plot(x='horsepower', y='mpg', style='o')
# plt.title("horsepower vs mpg")
# plt.ylabel('mpg')
# plt.plot(x, ypred)
# plt.show()


y_pred=reg.predict(x)
resid = y-y_pred

sm.graphics.influence_plot(model)
plt.show()

plt.plot(y_pred,resid,'o')
plt.title('Residuals vs Predicted Values')
plt.ylabel('Residual')
plt.xlabel('Prediction')
plt.show()

# Input vs Error
plt.plot(y,y_pred,'o')
plt.plot(y,y)
plt.title('Prediction vs True Value')
plt.xlabel('True Value')
plt.ylabel('Prediction')
plt.show()
```





The error vs predicted values chart seems to have a slight curve, which means that the relationship of the data is non-linearly associated. We might use non-linear transformations of the predictors such as $\log(x)$. The error terms also show a non-constant variance, which could be solved by using a non-linear transformation. The Influence plot reveals no real outliers, but reveals a very high leverage data point, point 13 (starting count from 0). Observation 28 also has a relatively high leverage. We can assess that observations 320, 324, 28 and 13 are likely outliers in our data along with other observations with a studentized residual of 3 or more. Most of our observations have a leverage lesser than 0.075 with the exception of two observations: 28, and 13.

e1)

```
In [ ]: model_interaction = smf.ols(formula='mpg ~ weight + cylinders + weight:cylinders', data=mpg)
summary = model_interaction.summary()
print(summary.tables[1])

model_interaction = smf.ols(formula='mpg ~ horsepower + acceleration + horsepower:acceleration', data=mpg)
summary = model_interaction.summary()
print(summary.tables[1])

model_interaction = smf.ols(formula='mpg ~ displacement + year + displacement:year', data=mpg)
summary = model_interaction.summary()
print(summary.tables[1])

model_interaction = smf.ols(formula='mpg ~ horsepower + displacement + horsepower:displacement', data=mpg)
summary = model_interaction.summary()
print(summary.tables[1])
```



```
summary = model_interaction.summary()
print(summary.tables[1])
```

```
=====
              coef      std err          t      P>|t|      [0.025      0.975]
-----
Intercept          65.3865       3.733      17.514      0.000       58.046       72.727
weight          -0.0128       0.001      -9.418      0.000       -0.016      -0.010
cylinders        -4.2098       0.724      -5.816      0.000       -5.633      -2.787
weight:cylinders    0.0011       0.000       5.226      0.000       0.001       0.002
=====
=====
              coef      std err          t      P>|t|      [0.025
0.975]
-----
Intercept          33.5124       3.420       9.798      0.000       26.788
40.237
horsepower          0.0176       0.027       0.641      0.522       -0.036
0.072
acceleration        0.8003       0.212       3.777      0.000       0.384
1.217
horsepower:acceleration -0.0157       0.002      -7.838      0.000      -0.020
-0.012
=====
=====
              coef      std err          t      P>|t|      [0.025      0.975]
-----
Intercept          -72.8784       8.368      -8.709      0.000      -89.330      -56.427
displacement         0.2523       0.041       6.216      0.000       0.173       0.332
year                 1.4077       0.110      12.779      0.000       1.191       1.624
displacement:year    -0.0041       0.001      -7.482      0.000      -0.005      -0.003
=====
=====
              coef      std err          t      P>|t|      [0.025
0.975]
-----
Intercept          53.0511       1.526      34.765      0.000       50.051
56.051
horsepower         -0.2343       0.020     -11.960      0.000      -0.273
-0.196
displacement        -0.0980       0.007     -14.674      0.000      -0.111
-0.085
horsepower:displacement 0.0006    5.19e-05     11.222      0.000       0.000
0.001
=====
=====
```

Interactions between weight and cylinders, and interactions between horsepower and acceleration seemed to be statistically significant (low p value). Interactions between displacement and year, and horsepower and displacement also seemed to be significant.

e2)

```
In [ ]: model_interaction = smf.ols(formula='mpg ~ weight:cylinders', data = df).fit()
summary = model_interaction.summary()
print(summary.tables[1])

model_interaction = smf.ols(formula='mpg ~ horsepower:acceleration', data = df).fit()
summary = model_interaction.summary()
print(summary.tables[1])

model_interaction = smf.ols(formula='mpg ~ displacement:weight', data = df).fit()
summary = model_interaction.summary()
print(summary.tables[1])
```

```
=====
              coef      std err          t      P>|t|      [0.025      0.975]
-----
Intercept          34.2948         0.464       73.906      0.000       33.382       35.207
weight:cylinders   -0.0006      2.28e-05     -27.029      0.000       -0.001       -0.001
=====
=====
              coef      std err          t      P>|t|      [0.025
0.975]
-----
Intercept          47.7299         0.978       48.789      0.000       45.807
49.653
horsepower:acceleration -0.0157         0.001     -25.612      0.000       -0.017
-0.014
=====
=====
              coef      std err          t      P>|t|      [0.025      0.97
5]
-----
--
Intercept          31.2634         0.388       80.588      0.000       30.501       32.0
26
displacement:weight -1.182e-05      4.6e-07     -25.687      0.000     -1.27e-05     -1.09e-
05
=====
==
```

There is a significance in the interaction between weight and cylinders, horsepower and acceleration, and displacement and weight.

f)

```
In [ ]: #Log
model = smf.ols(formula='mpg ~ cylinders + displacement + np.log(horsepower) + weight
summary = model.summary()
print("F value: {}".format(model.fvalue) )
print(model.summary().tables[1])
```

```
#quadratic
model = smf.ols(formula='mpg ~ cylinders + displacement + np.power(horsepower, 2) + weight + year + origin', data=mpg)
summary = model.summary()
print("F value: {}".format(model.fvalue) )
print(model.summary().tables[1])
```

F value: 286.8516549018523

```
=====
===
```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	42.9742	10.630	4.043	0.000	22.075	63.874
cylinders	-0.4384	0.305	-1.439	0.151	-1.037	0.161
displacement	0.0156	0.007	2.231	0.026	0.002	0.029
np.log(horsepower)	-10.4455	1.522	-6.863	0.000	-13.438	-7.453
weight	-0.0038	0.001	-5.359	0.000	-0.005	-0.002
np.log(acceleration)	-6.1583	1.645	-3.744	0.000	-9.392	-2.925
year	0.7050	0.048	14.683	0.000	0.611	0.799
origin	1.4379	0.258	5.574	0.000	0.931	1.945

```
=====
===
```

F value: 256.6597155187827

```
=====
=====
```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	-20.3992	4.117	-4.955	0.000	-28.494	-12.304
cylinders	-0.2820	0.327	-0.862	0.389	-0.925	0.361
displacement	0.0106	0.008	1.355	0.176	-0.005	0.026
np.power(horsepower, 2)	8.48e-05	4.14e-05	2.049	0.041	3.42e-06	0.000
weight	-0.0071	0.001	-12.208	0.000	-0.008	-0.006
np.power(acceleration, 2)	0.0080	0.003	3.194	0.002	0.003	0.013
year	0.7845	0.050	15.587	0.000	0.686	0.883
origin	1.1926	0.280	4.267	0.000	0.643	1.742

```
=====
=====
```

I tried performing a log transformation, and a quadratic transformation on displacement and

horsepower, as they seemed to have a non-linear relationship to mpg in the scatterplot matrix. The previous F value for the linear regression is 256.

The log transformation yielded an F value of 305, which means that we can reject the null hypothesis that the coefficients are equal to zero. The log transformation on displacement increased the p value of the coefficient, and on horsepower significantly decreased the p value of the coefficient. All of the log terms were significant with a 0.05 significance level.

The quadratic transformation yielded an F value of 256, so we can reject the null hypothesis that the coefficients are equal to zero. The quadratic transformation on displacement, horsepower, and weight decreased the p value on the displacement coefficient, slightly increased the p value on the horsepower coefficient. All of the squared terms were significant with a 0.05 significance level.

3.

a)

```
In [ ]: df = pd.read_csv("Carseats.csv")

#convert to 0 and 1
df['Urban'] = df['Urban'].eq('Yes').mul(1)
df['US'] = df['US'].eq('Yes').mul(1)

x = df[['Price', 'Urban', 'US']]

y = df['Sales'].values.reshape(-1, 1)
X = sm.add_constant(x)
model = sm.OLS(y,X).fit()

# model = smf.ols(formula='Sales ~ Price + Urban + US', data = df).fit()
# summary = model.summary()
print(model.summary())
print("RSE: {}".format(np.sqrt(model.scale)))
```

OLS Regression Results

=====					
Dep. Variable:	y	R-squared:	0.239		
Model:	OLS	Adj. R-squared:	0.234		
Method:	Least Squares	F-statistic:	41.52		
Date:	Tue, 15 Nov 2022	Prob (F-statistic):	2.39e-23		
Time:	18:40:09	Log-Likelihood:	-927.66		
No. Observations:	400	AIC:	1863.		
Df Residuals:	396	BIC:	1879.		
Df Model:	3				
Covariance Type:	nonrobust				
=====					
	coef	std err	t	P> t	[0.025 0.975]

const	13.0435	0.651	20.036	0.000	11.764 14.323
Price	-0.0545	0.005	-10.389	0.000	-0.065 -0.044
Urban	-0.0219	0.272	-0.081	0.936	-0.556 0.512
US	1.2006	0.259	4.635	0.000	0.691 1.710
=====					
Omnibus:	0.676	Durbin-Watson:	1.912		
Prob(Omnibus):	0.713	Jarque-Bera (JB):	0.758		
Skew:	0.093	Prob(JB):	0.684		
Kurtosis:	2.897	Cond. No.	628.		

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

RSE: 2.4724924402701642

b)

For each one unit increase in price, sales decrease by -0.0545 on average while US and Urban are constant. Sales are -0.0219 units lower for observations that are Urban while US and Price are constant, and 1.2006 units higher for observations that are US while Urban and Price are constant

c)

for non-urban, non-us: Sales = 13.0435 - 0.0545(Price)

for urban, non-us: Sales = 13.0435 - 0.0545(Price) - 0.0219 + e

for non-urban, us: Sales = 13.0435 - 0.0545(Price) + 1.2006 + e

for urban and us: Sales = 13.0435 - 0.0545(Price) - 0.0219 + 1.2006 + e

d)

we can reject the null hypothesis for predictors Price and US, as their p value is below our alpha of 0.05. We cannot reject the null hypothesis of Urban, as the p value is very high (above our alpha of 0.05).

e)

```
In [ ]: x = df[['Price', 'US']]
```

```

y = df['Sales'].values.reshape(-1, 1)
X = sm.add_constant(x)
model = sm.OLS(y,X).fit()

print(model.summary())
print("RSE: {}".format(np.sqrt(model.scale)))

```

```

                                OLS Regression Results
=====
Dep. Variable:                  y      R-squared:                0.239
Model:                        OLS      Adj. R-squared:           0.235
Method:                    Least Squares      F-statistic:            62.43
Date:                Tue, 15 Nov 2022      Prob (F-statistic):      2.66e-24
Time:                        18:40:09      Log-Likelihood:         -927.66
No. Observations:          400      AIC:                   1861.
Df Residuals:              397      BIC:                   1873.
Df Model:                    2
Covariance Type:            nonrobust
=====
               coef      std err          t      P>|t|      [0.025      0.975]
-----
const         13.0308         0.631     20.652     0.000     11.790     14.271
Price         -0.0545         0.005    -10.416     0.000     -0.065     -0.044
US             1.1996         0.258      4.641     0.000      0.692      1.708
=====
Omnibus:                 0.666      Durbin-Watson:           1.912
Prob(Omnibus):            0.717      Jarque-Bera (JB):         0.749
Skew:                     0.092      Prob(JB):                 0.688
Kurtosis:                 2.895      Cond. No.                  607.
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

RSE: 2.469396800574444

f) The models in a and e likely fit the data okay, as they have a low R-squared value of around 0.24 for both. The RSE of both models are above 1, but relatively close to 1. The model in e has a slightly better RSE, and R squared value, so it fits the data a bit better than the model in a.

g)

95% confidence interval

coefficient > 0.025 < 0.975

	coefficient	> 0.025	< 0.975
Price	-0.065		-0.044
US	0.692	1.708	

h)

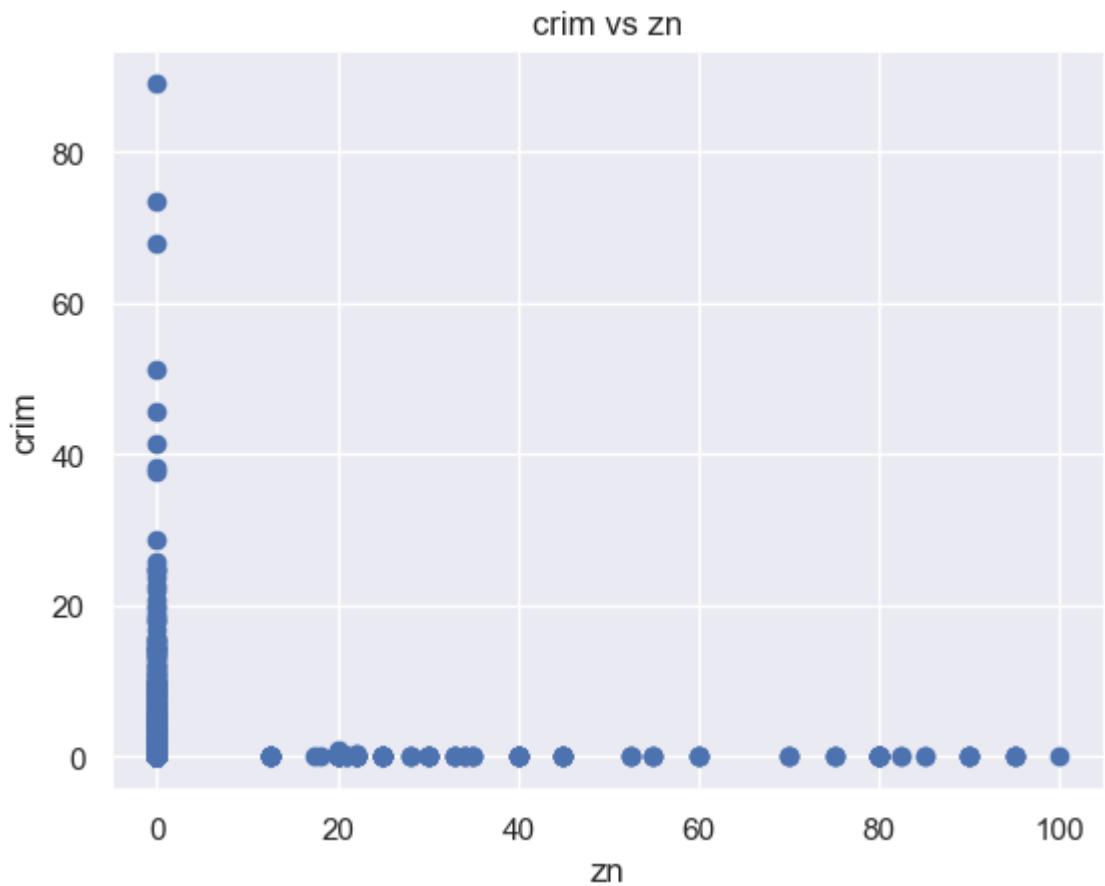
```

In [ ]: sm.graphics.influence_plot(model)
plt.show()

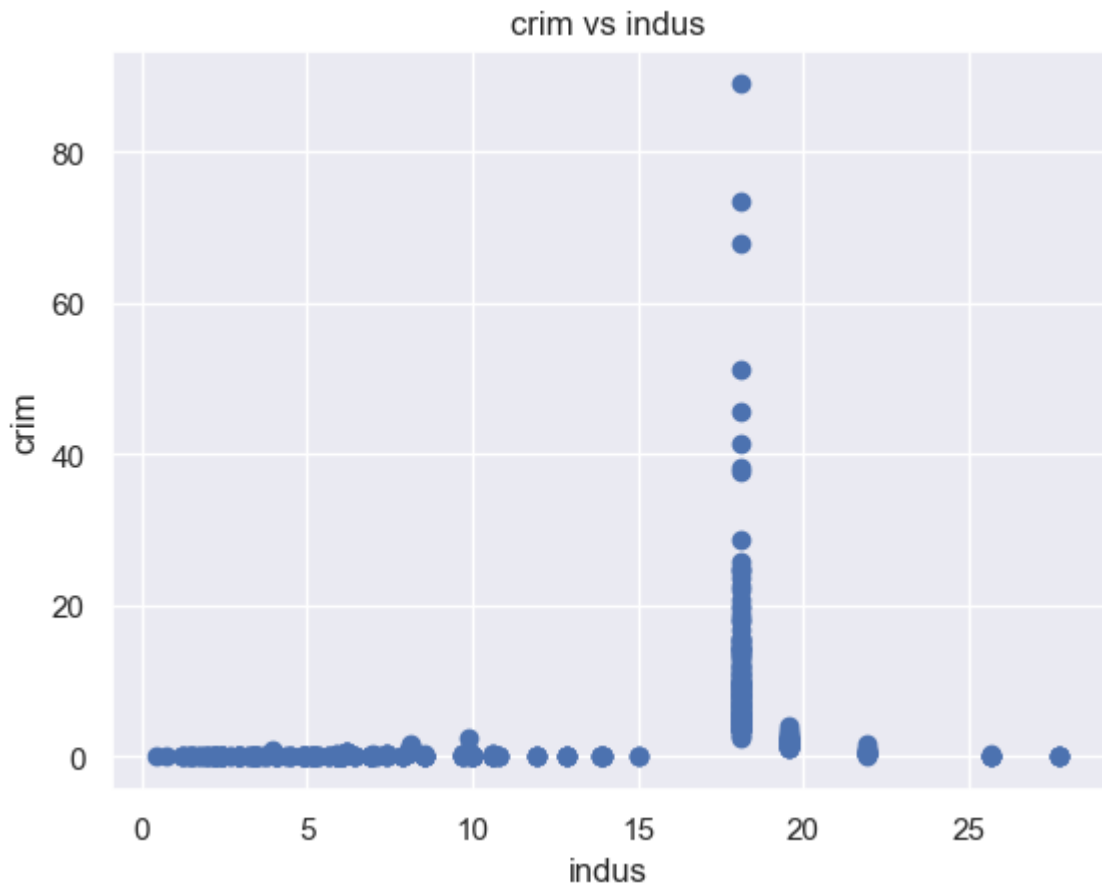
```



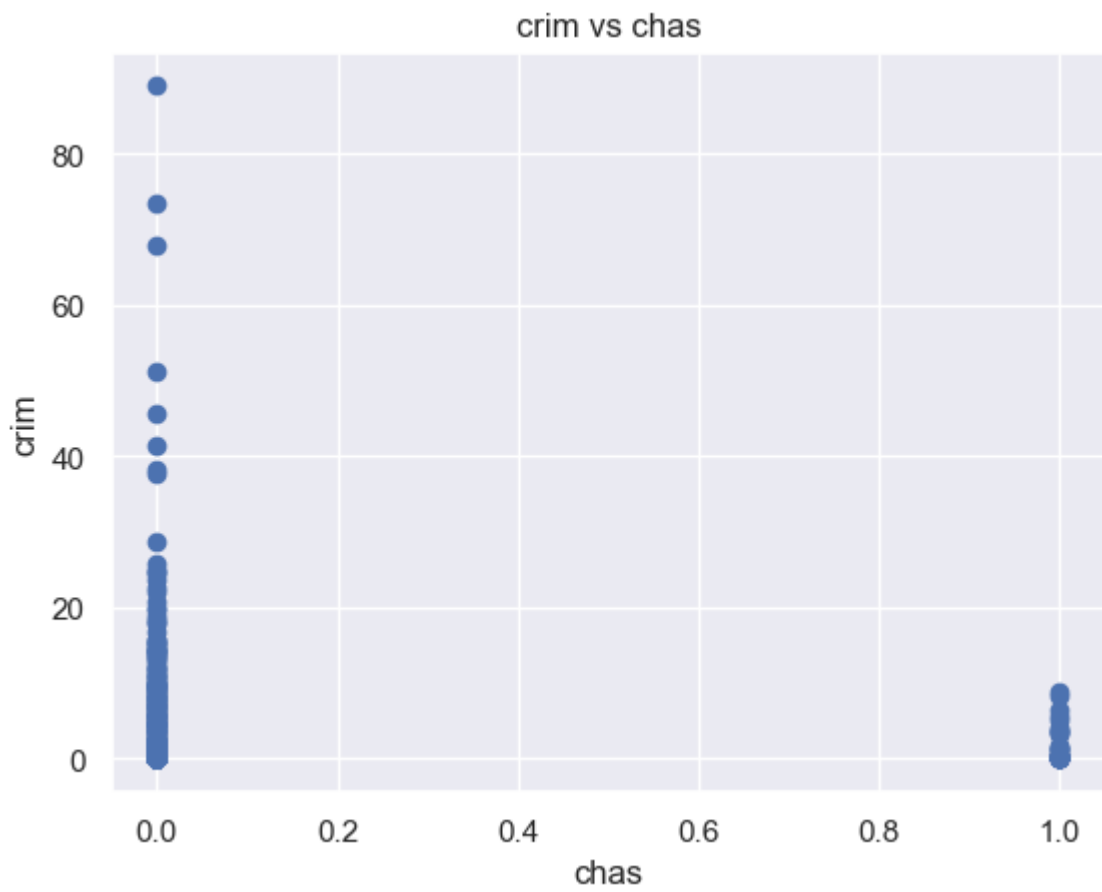
```
plt.show()
print(model.summary().tables[1])
univar_coefficients.append(model.params[1])
```




	coef	std err	t	P> t	[0.025	0.975]
const	4.4537	0.417	10.675	0.000	3.634	5.273
zn	-0.0739	0.016	-4.594	0.000	-0.106	-0.042



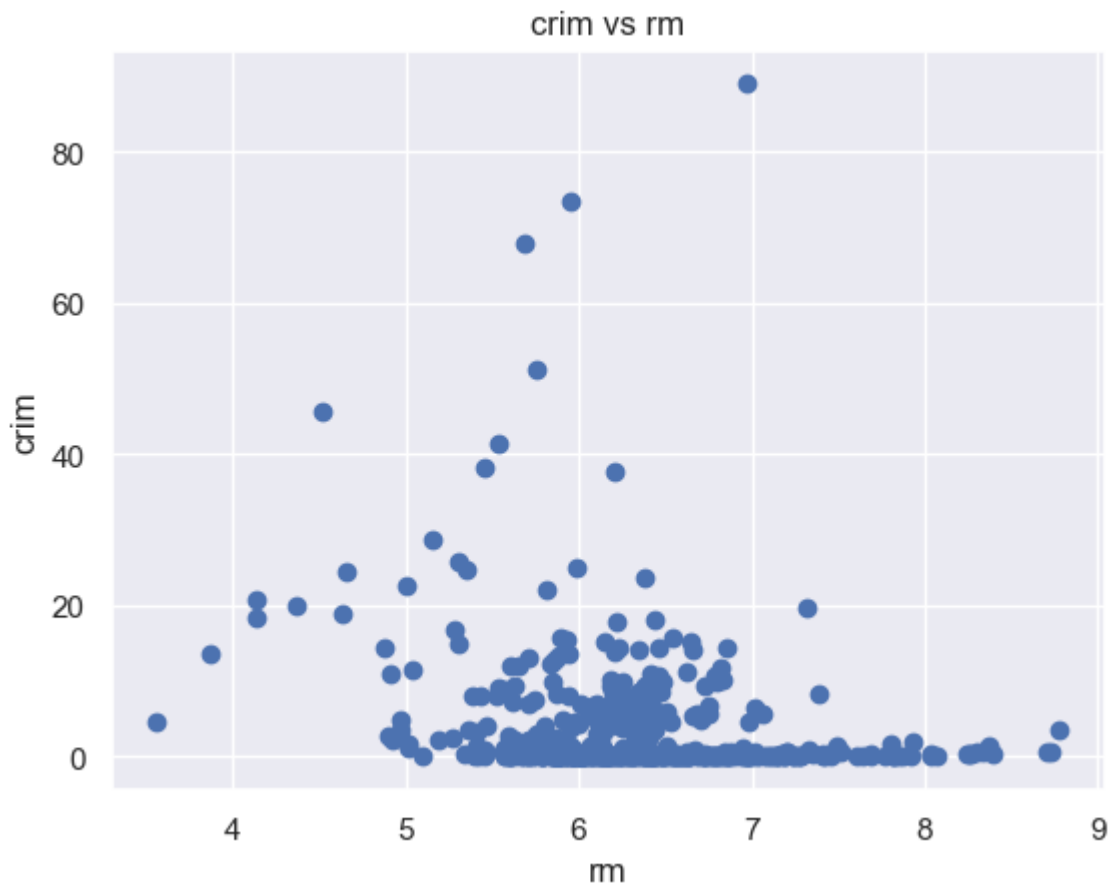
	coef	std err	t	P> t	[0.025	0.975]
const	-2.0637	0.667	-3.093	0.002	-3.375	-0.753
indus	0.5098	0.051	9.991	0.000	0.410	0.610



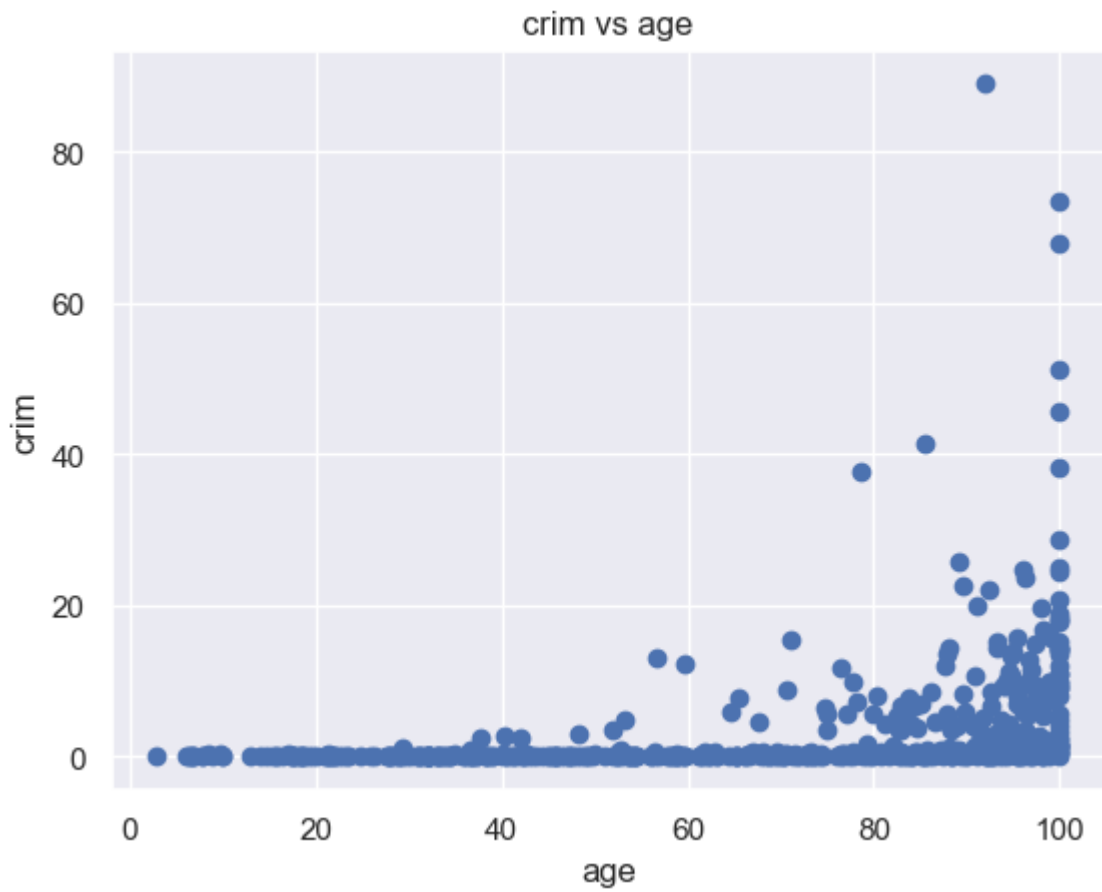
	coef	std err	t	P> t	[0.025	0.975]
const	3.7444	0.396	9.453	0.000	2.966	4.523
chas	-1.8928	1.506	-1.257	0.209	-4.852	1.066




	coef	std err	t	P> t	[0.025	0.975]
const	-13.7199	1.699	-8.073	0.000	-17.059	-10.381
nox	31.2485	2.999	10.419	0.000	25.356	37.141



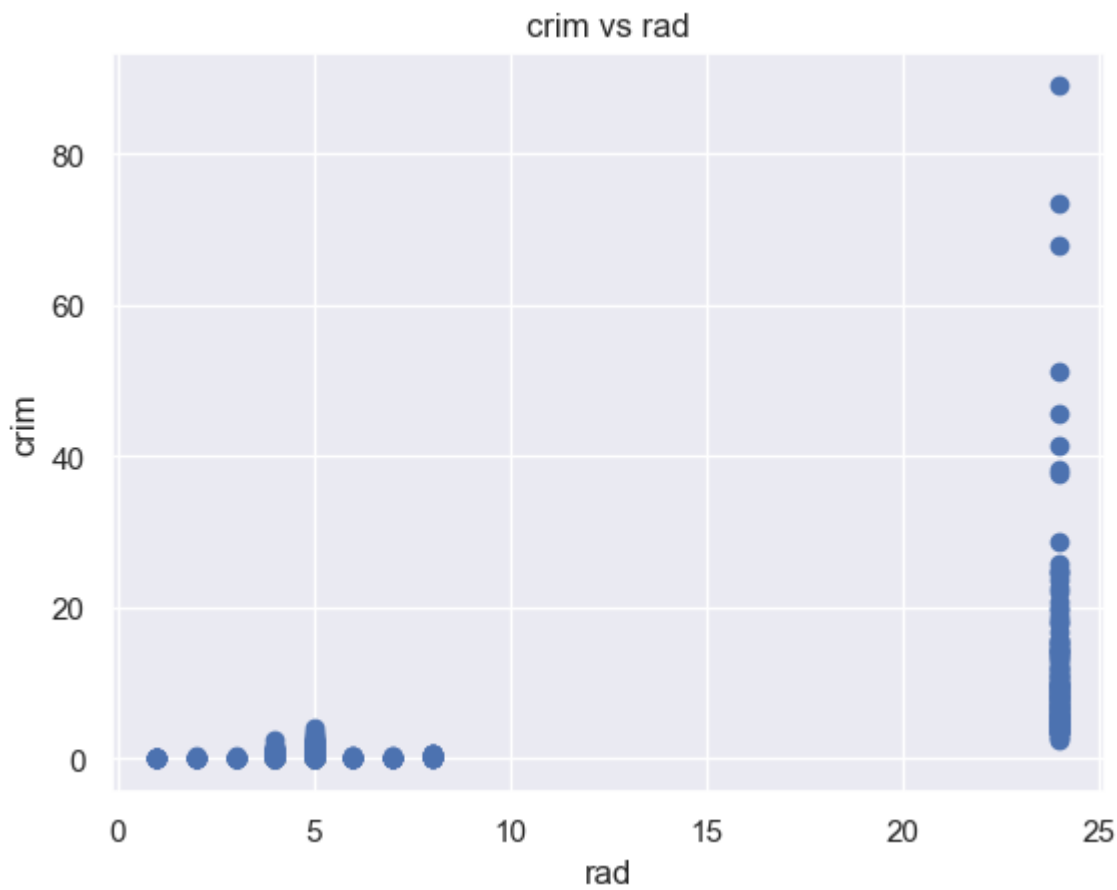
	coef	std err	t	P> t	[0.025	0.975]
const	20.4818	3.364	6.088	0.000	13.872	27.092
rm	-2.6841	0.532	-5.045	0.000	-3.729	-1.639



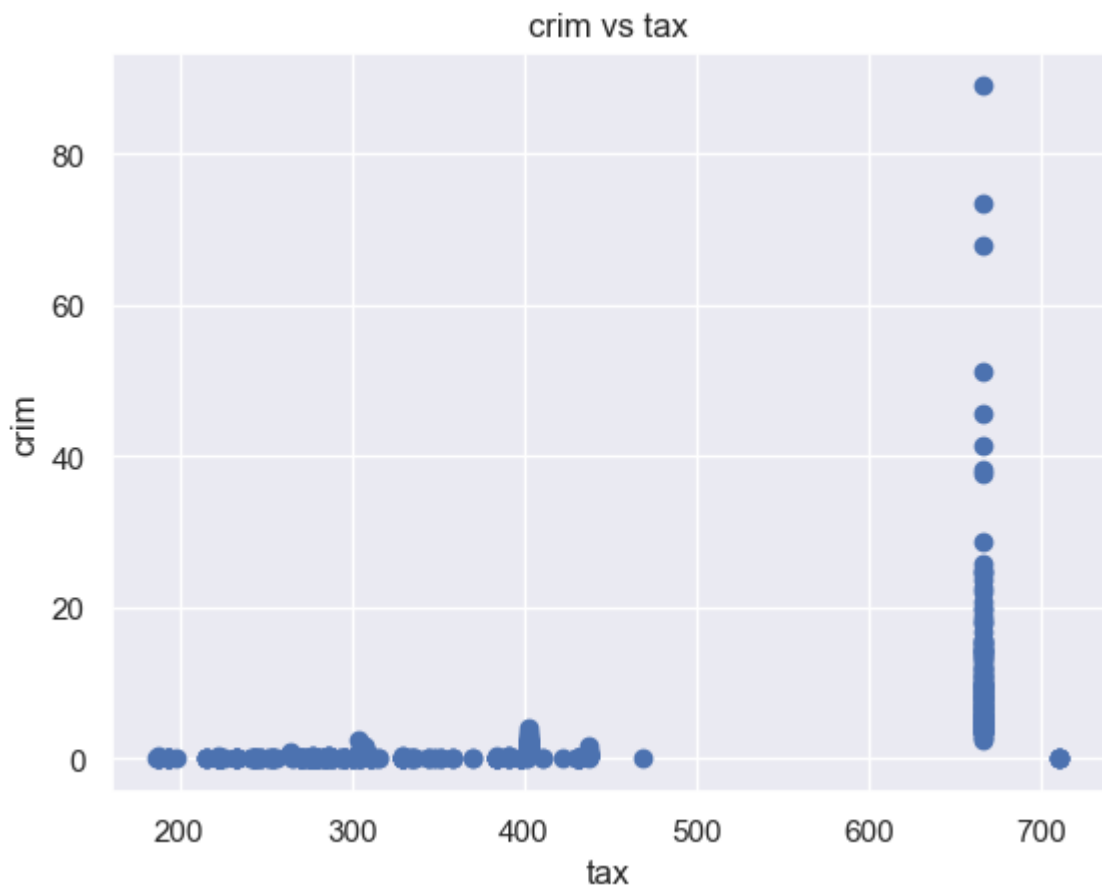
	coef	std err	t	P> t	[0.025	0.975]
const	-3.7779	0.944	-4.002	0.000	-5.633	-1.923
age	0.1078	0.013	8.463	0.000	0.083	0.133



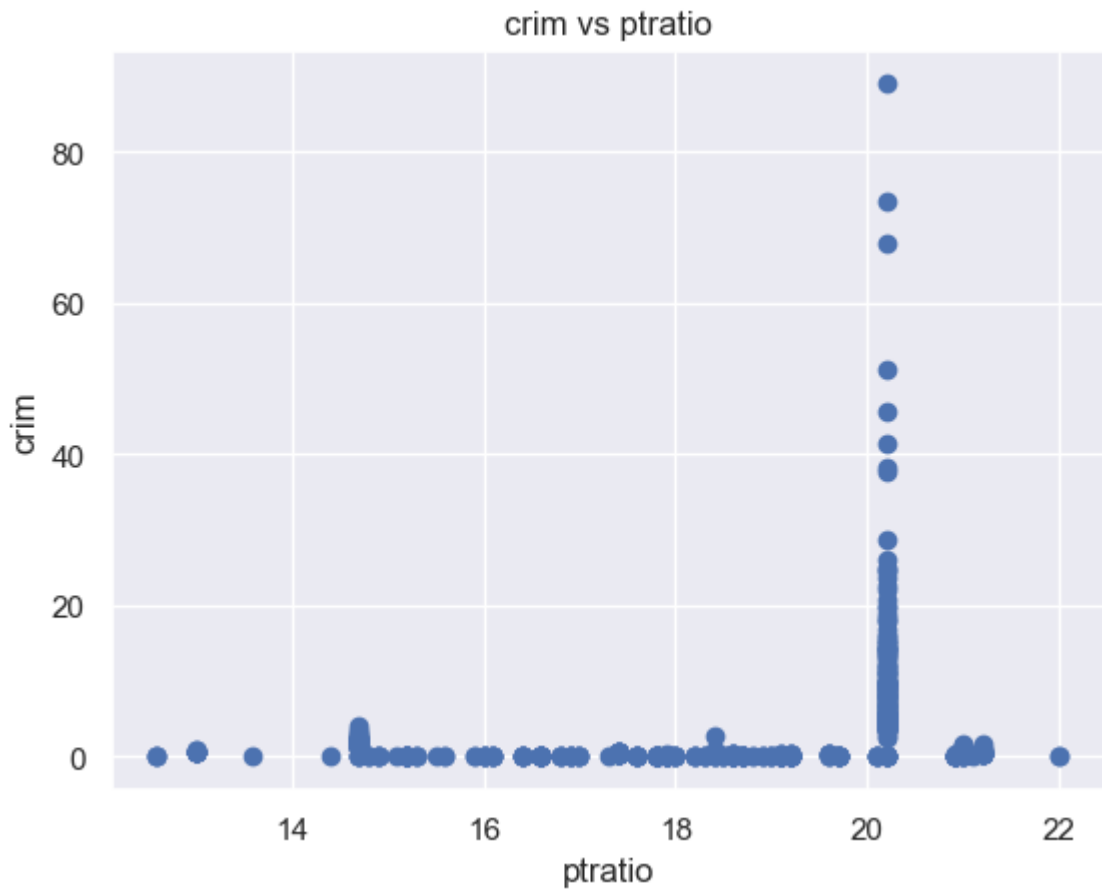
	coef	std err	t	P> t	[0.025	0.975]
const	9.4993	0.730	13.006	0.000	8.064	10.934
dis	-1.5509	0.168	-9.213	0.000	-1.882	-1.220



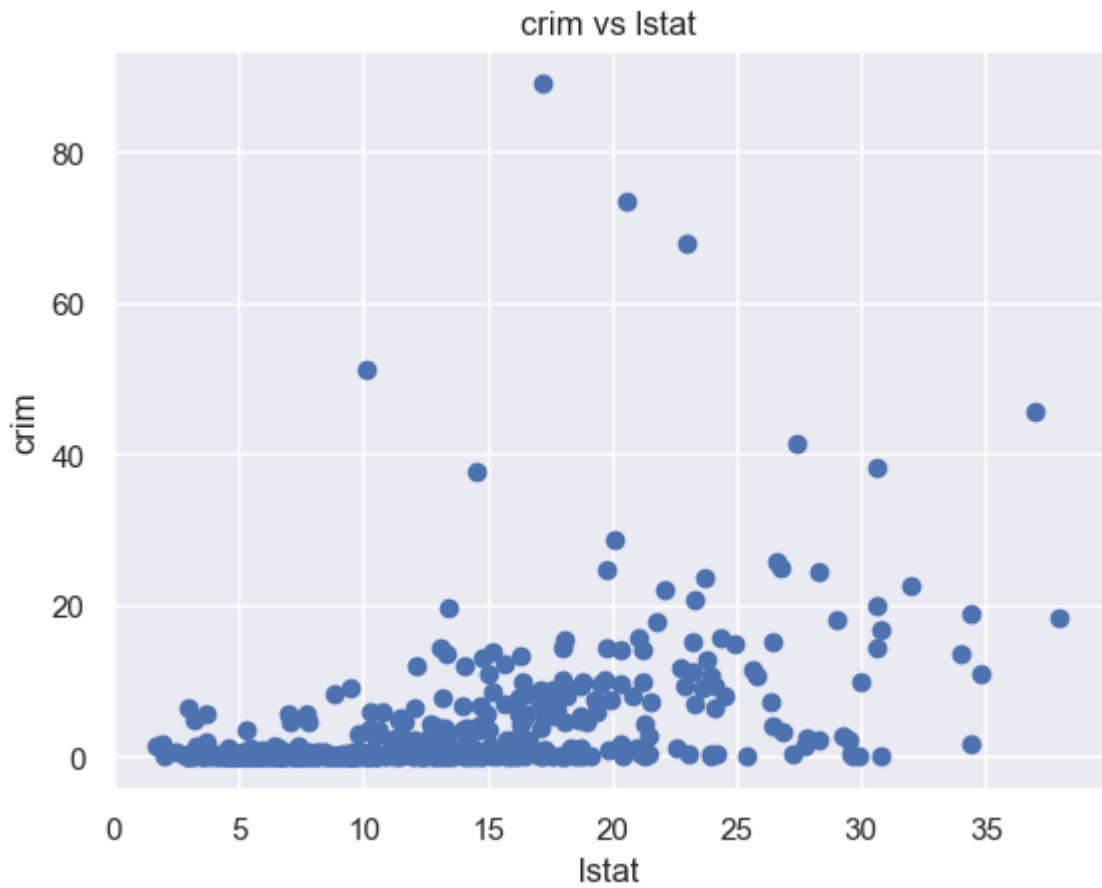
	coef	std err	t	P> t	[0.025	0.975]
const	-2.2872	0.443	-5.157	0.000	-3.158	-1.416
rad	0.6179	0.034	17.998	0.000	0.550	0.685




	coef	std err	t	P> t	[0.025	0.975]
const	-8.5284	0.816	-10.454	0.000	-10.131	-6.926
tax	0.0297	0.002	16.099	0.000	0.026	0.033



	coef	std err	t	P> t	[0.025	0.975]
const	-17.6469	3.147	-5.607	0.000	-23.830	-11.464
ptratio	1.1520	0.169	6.801	0.000	0.819	1.485



	coef	std err	t	P> t	[0.025	0.975]
const	-3.3305	0.694	-4.801	0.000	-4.694	-1.968
lstat	0.5488	0.048	11.491	0.000	0.455	0.643



	coef	std err	t	P> t	[0.025	0.975]
const	11.7965	0.934	12.628	0.000	9.961	13.632
medv	-0.3632	0.038	-9.460	0.000	-0.439	-0.288

all predictors have a p value of < 0.05 except for 'chas' so we acn determine that there is a statistically significant association between the predictors and response except for in 'chas'.

b)

```
In [ ]: X = sm.add_constant(all_x)
model = sm.OLS(y,X).fit()

print(model.summary())
```

OLS Regression Results

=====						
Dep. Variable:	y	R-squared:	0.449			
Model:	OLS	Adj. R-squared:	0.436			
Method:	Least Squares	F-statistic:	33.52			
Date:	Tue, 15 Nov 2022	Prob (F-statistic):	2.03e-56			
Time:	18:40:11	Log-Likelihood:	-1655.4			
No. Observations:	506	AIC:	3337.			
Df Residuals:	493	BIC:	3392.			
Df Model:	12					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]

const	13.7784	7.082	1.946	0.052	-0.136	27.693
zn	0.0457	0.019	2.433	0.015	0.009	0.083
indus	-0.0584	0.084	-0.698	0.486	-0.223	0.106
chas	-0.8254	1.183	-0.697	0.486	-3.150	1.500
nox	-9.9576	5.290	-1.882	0.060	-20.351	0.436
rm	0.6289	0.607	1.036	0.301	-0.564	1.822
age	-0.0008	0.018	-0.047	0.962	-0.036	0.034
dis	-1.0122	0.282	-3.584	0.000	-1.567	-0.457
rad	0.6125	0.088	6.997	0.000	0.440	0.784
tax	-0.0038	0.005	-0.730	0.466	-0.014	0.006
ptratio	-0.3041	0.186	-1.632	0.103	-0.670	0.062
lstat	0.1388	0.076	1.833	0.067	-0.010	0.288
medv	-0.2201	0.060	-3.678	0.000	-0.338	-0.103
=====						
Omnibus:	663.436	Durbin-Watson:	1.516			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	80856.852			
Skew:	6.579	Prob(JB):	0.00			
Kurtosis:	63.514	Cond. No.	1.24e+04			
=====						

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.24e+04. This might indicate that there are strong multicollinearity or other numerical problems.

The r squared is relatively high, which means that the model is a relatively good fit for the data.


We can reject the null hypothesis for zn, dis, rad, and medv, as their p values are below our alpha of 0.05.

c)

```
In [ ]: multiple_reg = model.params[1:]
data = pd.DataFrame(multiple_reg, columns = ['multiple'] )
data['univar'] = univar_coefficients
plt.scatter(data['univar'], data['multiple'])
plt.xlabel('simple')
plt.ylabel('multiple')

for index, row in data.iterrows():
    plt.annotate(index, (row['univar'], row['multiple']))
```

```
plt.show()
```



The results showed that there is a difference for the simple and multiple regression coefficients. This is due to the fact that in the simple regression, the slope represents the average effect of the predictor, ignoring other factors. In a multiple regression, the slope represents the average effect of the predictor holding other factors constant.

d)

```
In [ ]: #we want to predict per capita crime rate.
df = pd.read_csv("Boston.csv", index_col=0)
y = df['crim'].values.reshape(-1, 1)
all_x = df.drop(['crim'], axis=1)

univar_coefficients = []
# y = df['Sales'].values.reshape(-1, 1)
for col in all_x:
    x = df[col]
    model = smf.ols(formula='crim ~ {} + np.power({}, 2) + np.power({}, 3)'.format(col), data=df)
    print(model.summary().tables[1])
```

=====						
	coef	std err	t	P> t	[0.025	0.975]

Intercept	4.8461	0.433	11.192	0.000	3.995	5.697
zn	-0.3322	0.110	-3.025	0.003	-0.548	-0.116
np.power(zn, 2)	0.0065	0.004	1.679	0.094	-0.001	0.014
np.power(zn, 3)	-3.776e-05	3.14e-05	-1.203	0.230	-9.94e-05	2.39e-05
=====						
=						
	coef	std err	t	P> t	[0.025	0.97

5]						

-						
Intercept	3.6626	1.574	2.327	0.020	0.570	6.75
indus	-1.9652	0.482	-4.077	0.000	-2.912	-1.01
np.power(indus, 2)	0.2519	0.039	6.407	0.000	0.175	0.32
np.power(indus, 3)	-0.0070	0.001	-7.292	0.000	-0.009	-0.00
=====						
=						

	coef	std err	t	P> t	[0.025	0.975]

Intercept	3.7444	0.397	9.444	0.000	2.965	4.523
chas	1.114e+14	2.71e+14	0.411	0.681	-4.21e+14	6.44e+14
np.power(chas, 2)	-5.61e+13	1.37e+14	-0.411	0.681	-3.24e+14	2.12e+14
np.power(chas, 3)	-5.532e+13	1.35e+14	-0.411	0.681	-3.2e+14	2.09e+14
=====						

	coef	std err	t	P> t	[0.025	0.975]

Intercept	233.0866	33.643	6.928	0.000	166.988	299.185
nox	-1279.3713	170.397	-7.508	0.000	-1614.151	-944.591
np.power(nox, 2)	2248.5441	279.899	8.033	0.000	1698.626	2798.462
np.power(nox, 3)	-1245.7029	149.282	-8.345	0.000	-1538.997	-952.409
=====						

	coef	std err	t	P> t	[0.025	0.975]

Intercept	112.6246	64.517	1.746	0.081	-14.132	239.382
rm	-39.1501	31.311	-1.250	0.212	-100.668	22.368
np.power(rm, 2)	4.5509	5.010	0.908	0.364	-5.292	14.394
np.power(rm, 3)	-0.1745	0.264	-0.662	0.509	-0.693	0.344
=====						

	coef	std err	t	P> t	[0.025	0.975]

Intercept	-2.5488	2.769	-0.920	0.358	-7.989	2.892
age	0.2737	0.186	1.468	0.143	-0.093	0.640
np.power(age, 2)	-0.0072	0.004	-1.988	0.047	-0.014	-8.4e-05
np.power(age, 3)	5.745e-05	2.11e-05	2.724	0.007	1.6e-05	9.89e-05
=====						
=====						

	coef	std err	t	P> t	[0.025	0.975]

Intercept	30.0476	2.446	12.285	0.000	25.242	34.853
dis	-15.5544	1.736	-8.960	0.000	-18.965	-12.144
np.power(dis, 2)	2.4521	0.346	7.078	0.000	1.771	3.133
np.power(dis, 3)	-0.1186	0.020	-5.814	0.000	-0.159	-0.079
=====						
=====						
	coef	std err	t	P> t	[0.025	0.975]

Intercept	-0.6055	2.050	-0.295	0.768	-4.633	3.422
rad	0.5127	1.044	0.491	0.623	-1.538	2.563
np.power(rad, 2)	-0.0752	0.149	-0.506	0.613	-0.367	0.217
np.power(rad, 3)	0.0032	0.005	0.703	0.482	-0.006	0.012
=====						
=====						
	coef	std err	t	P> t	[0.025	0.975]

Intercept	19.1836	11.796	1.626	0.105	-3.991	42.358
tax	-0.1533	0.096	-1.602	0.110	-0.341	0.035
np.power(tax, 2)	0.0004	0.000	1.488	0.137	-0.000	0.001
np.power(tax, 3)	-2.204e-07	1.89e-07	-1.167	0.244	-5.91e-07	1.51e-07
=====						
=====						
	coef	std err	t	P> t	[0.025	0.975]

Intercept	477.1840	156.795	3.043	0.002	169.129	785.239
ptratio	-82.3605	27.644	-2.979	0.003	-136.673	-28.048
np.power(ptratio, 2)	4.6353	1.608	2.882	0.004	1.475	7.795
np.power(ptratio, 3)	-0.0848	0.031	-2.743	0.006	-0.145	-0.024
=====						
=====						
	coef	std err	t	P> t	[0.025	0.975]

Intercept	1.2010	2.029	0.592	0.554	-2.785	5.187
lstat	-0.4491	0.465	-0.966	0.335	-1.362	0.464
np.power(lstat, 2)	0.0558	0.030	1.852	0.065	-0.003	0.115
np.power(lstat, 3)	-0.0009	0.001	-1.517	0.130	-0.002	0.000
=====						
=====						
	coef	std err	t	P> t	[0.025	0.975]

Intercept	53.1655	3.356	15.840	0.000	46.571	59.760
medv	-5.0948	0.434	-11.744	0.000	-5.947	-4.242
np.power(medv, 2)	0.1555	0.017	9.046	0.000	0.122	0.189
np.power(medv, 3)	-0.0015	0.000	-7.312	0.000	-0.002	-0.001

=====

For zn, chas, rm, rad, tax, and lstat there was no statistical significance pointing to a non-linear association. For indus, nox, age, dis, ptratio, and medv there was a statistical significance pointing to a non-linear association.